

Prediction of Liver Disease within Individuals

Load the dataset

```
In [1]: import pandas as pd


file_path = 'Indian_Liver_Patients.csv'
data = pd.read_csv(file_path)
```

Explore the dataset

```
In [2]: data.head()
```

```
Out[2]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotran:
0	65	Female	0.7	0.1	187	
1	62	Male	10.9	5.5	699	
2	62	Male	7.3	4.1	490	
3	58	Male	1.0	0.4	182	
4	72	Male	3.9	2.0	195	



```
In [3]: #Get number of records and columns
data.shape
```

```
Out[3]: (583, 11)
```

```
In [4]: #Get data types
data.dtypes
```

```
Out[4]: Age                int64
Gender                object
Total_Bilirubin      float64
Direct_Bilirubin     float64
Alkaline_Phosphotase  int64
Alamine_Aminotransferase  int64
Aspartate_Aminotransferase  int64
Total_Protiens       float64
Albumin              float64
Albumin_and_Globulin_Ratio  float64
Dataset              int64
dtype: object
```

```
In [5]: #Get number of unique values in each column  
data.nunique()
```

```
Out[5]: Age                72  
        Gender             2  
        Total_Bilirubin    113  
        Direct_Bilirubin   80  
        Alkaline_Phosphotase 263  
        Alamine_Aminotransferase 152  
        Aspartate_Aminotransferase 177  
        Total_Protiens      58  
        Albumin            40  
        Albumin_and_Globulin_Ratio 69  
        Dataset            2  
        dtype: int64
```

Perform column conversions

```
In [6]: #Rename the column 'Dataset' to 'Target' for identification convenience  
data.rename(columns = {'Dataset':'Target'}, inplace = True)  
  
#Replace values in the 'Target' column: 2 -> 0  
data['Target'].replace({2: 0}, inplace = True)  
  
data['Target'].info()
```

```
<class 'pandas.core.series.Series'>  
RangeIndex: 583 entries, 0 to 582  
Series name: Target  
Non-Null Count  Dtype  
-----  
583 non-null    int64  
dtypes: int64(1)  
memory usage: 4.7 KB
```

```
In [7]: #Get counts on the 'Target' variable  
data['Target'].value_counts()
```

```
Out[7]: Target  
1      416  
0      167  
Name: count, dtype: int64
```

Handle duplicates

```
In [8]: #Check for duplicated count of records  
data.duplicated().sum()
```

```
Out[8]: 13
```

```
In [9]: #Drop them  
data = data.drop_duplicates()
```

```
data.shape[0]
```

Out[9]: 570

Handle missing values

```
In [10]: #Check for missing values
data.isnull().sum()
```

```
Out[10]: Age                                0
Gender                                      0
Total_Bilirubin                           0
Direct_Bilirubin                          0
Alkaline_Phosphotase                      0
Alamine_Aminotransferase                  0
Aspartate_Aminotransferase                0
Total_Protiens                            0
Albumin                                   0
Albumin_and_Globulin_Ratio                4
Target                                    0
dtype: int64
```

```
In [11]: #Impute the Albumin_and_Globulin_Ratio column with mean
data['Albumin_and_Globulin_Ratio'].fillna(data['Albumin_and_Globulin_Ratio'].mean())
#Verification
data.isnull().sum().sum()
```

Out[11]: 0

Encode categorical variables

```
In [12]: ## 'Gender'
data['Gender'] = data['Gender'].map({'Female':0, 'Male':1})
#Verification
print(data['Gender'].unique())
print(data.dtypes)
```

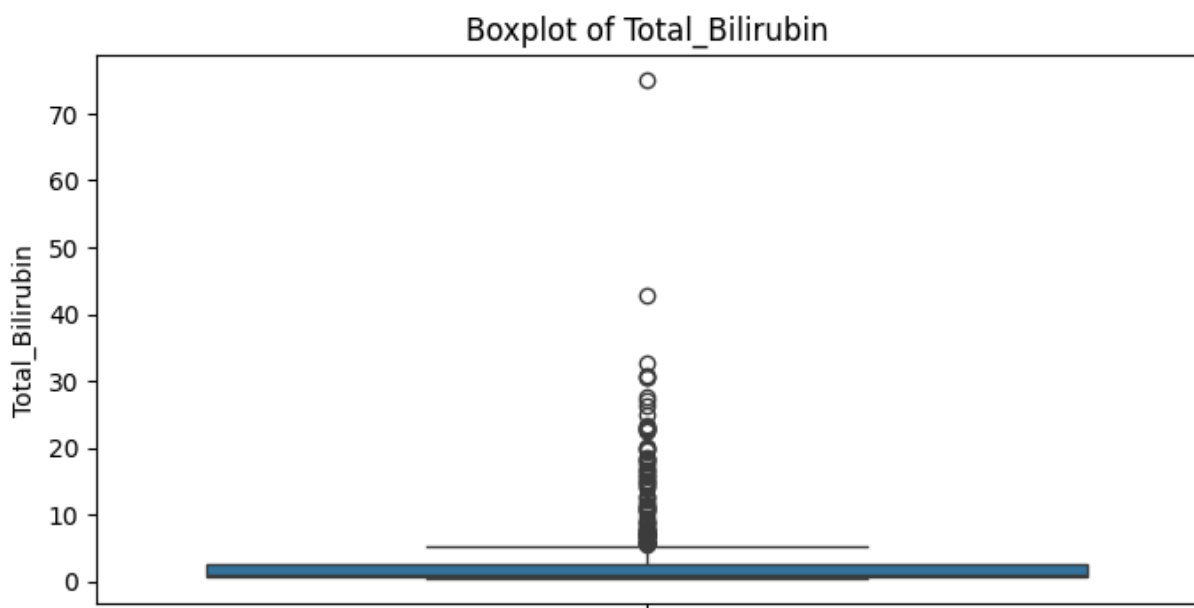

```
[0 1]
Age                                int64
Gender                            int64
Total_Bilirubin                   float64
Direct_Bilirubin                  float64
Alkaline_Phosphotase              int64
Alamine_Aminotransferase          int64
Aspartate_Aminotransferase        int64
Total_Protiens                    float64
Albumin                           float64
Albumin_and_Globulin_Ratio        float64
Target                            int64
dtype: object
```

Check for outliers through boxplots


```
In [13]: numerical_cols = data.drop(columns = ['Gender', 'Target'])
```

```
In [14]: import matplotlib.pyplot as plt
import seaborn as sns


for col in numerical_cols:
    plt.figure(figsize = (8, 4))
    sns.boxplot(data[col])
    plt.title(f'Boxplot of {col}')
    plt.show()
```




Boxplot of Direct_Bilirubin




Boxplot of Alkaline_Phosphotase




Boxplot of Alamine_Aminotransferase




Boxplot of Aspartate_Aminotransferase

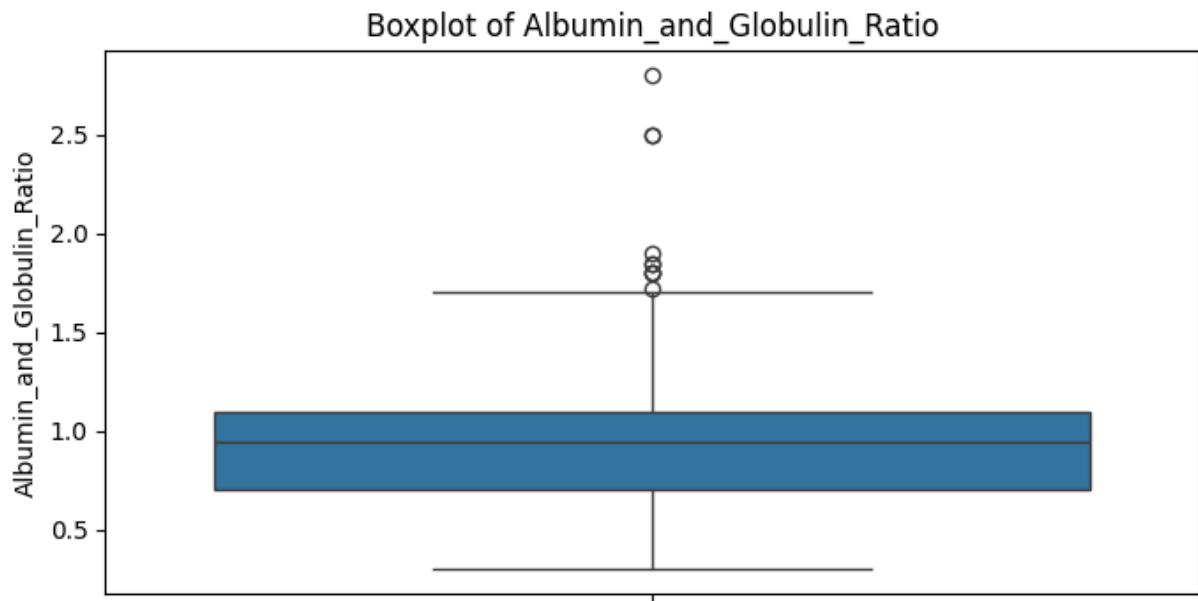


Boxplot of Total_Protiens



Boxplot of Albumin





```
In [15]: #Detect considerable amount of outliers  
#Perform robust scaling
```


Numerical data - Robust scaling

```
In [16]: from sklearn.preprocessing import RobustScaler  
  
numerical_names = data.drop(columns = ['Gender', 'Target']).columns  
  
scaler = RobustScaler()  
data[numerical_names] = scaler.fit_transform(data[numerical_names])  
data[numerical_names] = pd.DataFrame(data[numerical_names], columns = numerical_names)  
data
```

Out[16]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminot
0	0.80	0	-0.166667	-0.181818	-0.172131	
1	0.68	1	5.500000	4.727273	4.024590	
2	0.68	1	3.500000	3.454545	2.311475	
3	0.52	1	0.000000	0.090909	-0.213115	
4	1.08	1	1.611111	1.545455	-0.106557	
...
578	0.60	1	-0.277778	-0.181818	2.393443	
579	-0.20	1	-0.222222	-0.181818	-0.901639	
580	0.28	1	-0.111111	-0.090909	0.303279	
581	-0.56	1	0.166667	0.181818	-0.196721	
582	-0.28	1	0.000000	0.000000	0.065574	


570 rows × 11 columns



Distribution of existence of liver disease


```
In [17]: plt.figure(figsize = (6, 4))
sns.countplot(data = data, x = 'Target')
plt.title('Distribution of Response Variable')
plt.show()

#Detect heavy class imbalance
```

Correlation matrix

```
In [18]: plt.figure(figsize = (12, 8))
correlation_matrix = data.corr()
sns.heatmap(correlation_matrix, annot = True, cmap = 'coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



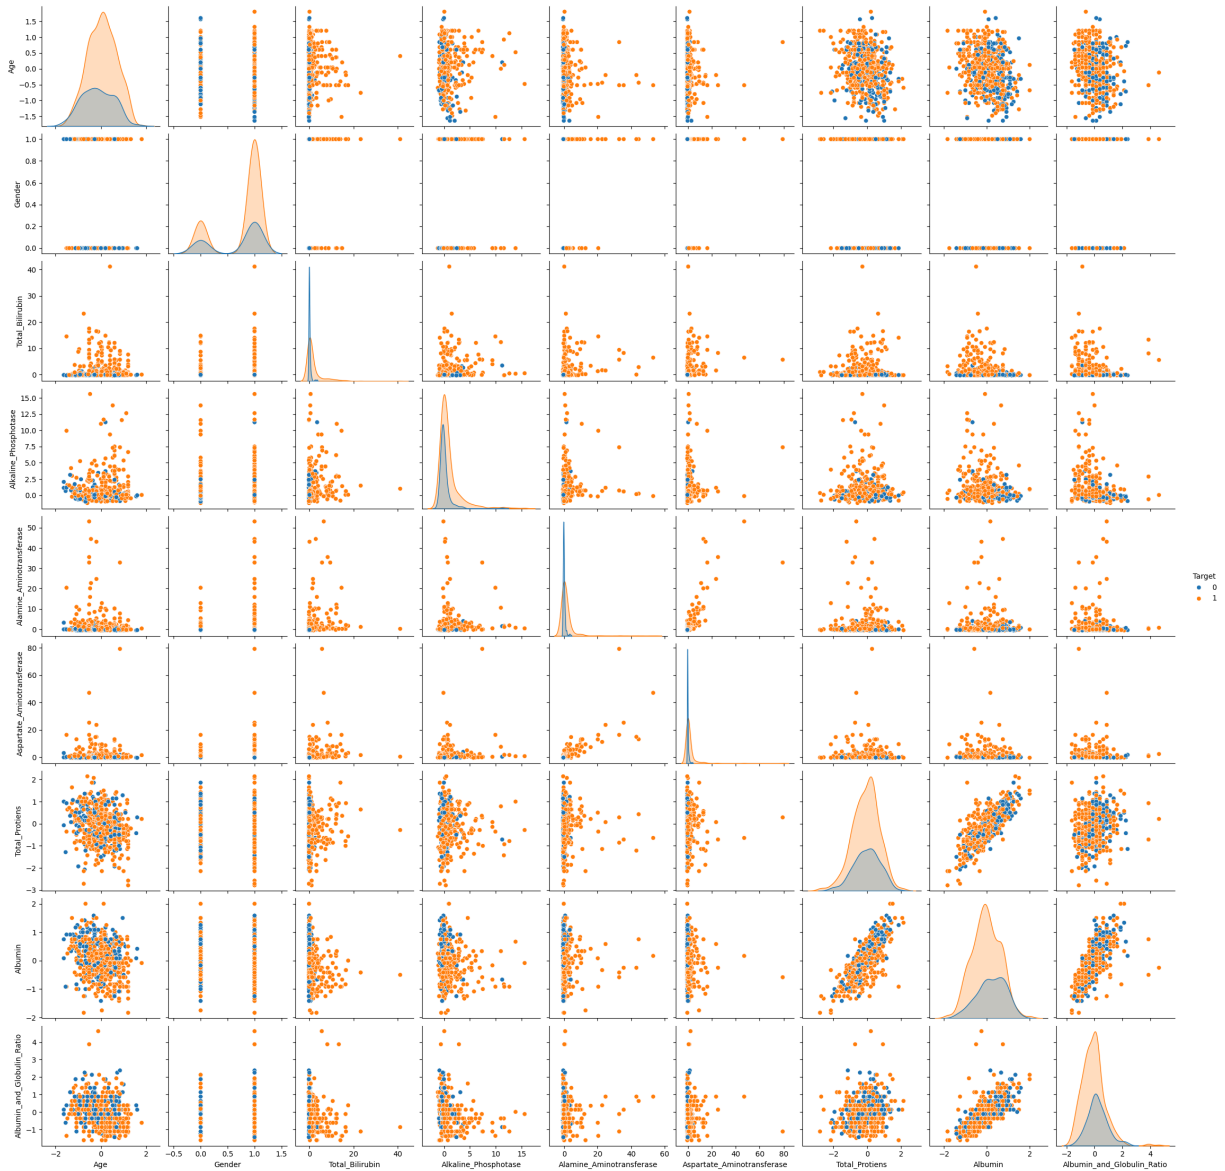
Drop one of the highly correlated columns

```
In [19]: data = data.drop(columns = ['Direct_Bilirubin'])
data.shape
```

```
Out[19]: (570, 10)
```

Pair plot distributions

```
In [20]: sns.pairplot(data, hue = 'Target', diag_kind = 'kde')
plt.show()
```



Split the data into train and test

```
In [21]: from sklearn.model_selection import train_test_split
from collections import Counter

X = data.drop('Target', axis = 1)
y = data['Target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_s
print("Original class distribution:", Counter(y_train))
```

Original class distribution: Counter({1: 321, 0: 135})

Perform SMOTE on response variable

```
In [22]: from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state = 42)
```

```
X_train, y_train = smote.fit_resample(X_train, y_train)
print("SMOTE class distribution:", Counter(y_train))
```

SMOTE class distribution: Counter({1: 321, 0: 321})

In [23]: X_train.shape

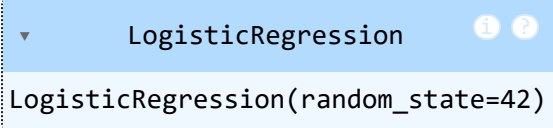
Out[23]: (642, 9)

Train models

1. Logistic Regression model

In [24]: `from` sklearn.linear_model `import` LogisticRegression

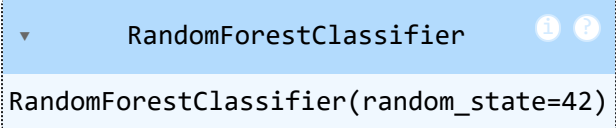
```
lr = LogisticRegression(random_state = 42, penalty = 'l2')
lr.fit(X_train, y_train)
```

Out[24]:  LogisticRegression
LogisticRegression(random_state=42)

2. Random Forest Classifier model

In [25]: `from` sklearn.ensemble `import` RandomForestClassifier

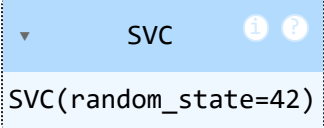
```
rf = RandomForestClassifier(random_state = 42)
rf.fit(X_train, y_train)
```

Out[25]:  RandomForestClassifier
RandomForestClassifier(random_state=42)

3. Support Vector Machine (SVM) model

In [26]: `from` sklearn.svm `import` SVC

```
svm = SVC(random_state = 42)
svm.fit(X_train, y_train)
```

Out[26]:  SVC
SVC(random_state=42)

4. K Nearest Neighbors (KNN) Classifier Model

```
In [27]: from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors = 2)
knn.fit(X_train, y_train)
```

```
Out[27]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=2)
```

5. XGBoost Classifier Model

```
In [28]: from xgboost import XGBClassifier
```

```
xgb = XGBClassifier(random_state = 42)
xgb.fit(X_train, y_train)
```

```
Out[28]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds
              =None,
              enable_categorical=False, eval_metric=None, feature_types
              =None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin
              =None,
```

Hyperparameter tuning & Cross Validation for models

```
In [29]: from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import StratifiedKFold

cv = StratifiedKFold(n_splits = 10, shuffle = True, random_state = 42)
```

1. Logistic Regression model

```
In [30]: #Define hyperparameter grids
param_grid_lr = {
    'C': [0.01, 0.1, 1, 10], #Regularization strength; inversely proportional to C
    'solver': ['newton-cg', 'lbfgs', 'liblinear'], #Algorithm used in optimization
}

#Perform GridSearchCV
grid_search_lr = GridSearchCV(lr, param_grid_lr, cv = 5, n_jobs = -1, scoring = 'ac
grid_search_lr.fit(X_train, y_train)
```

```
best_lr = grid_search_lr.best_estimator_  
best_lr_params = grid_search_lr.best_params_  
print("LR best parameters: ", best_lr_params)
```

LR best parameters: {'C': 0.1, 'solver': 'newton-cg'}

2. Random Forest Classifier model

```
In [31]: #Define hyperparameter grids  
param_grid_rf = {  
    'n_estimators': [500, 600, 700], #No of trees in the forest  
    'max_depth': [None, 10, 20], #Maximum depth of the tree  
    #None means nodes are expanded until all leaves are pure  
    'min_samples_split': [0.5, 1.0, 2, 5], #Minimum number of samples required to s  
    'min_samples_leaf': [1, 2, 4], #Minimum number of samples required to be at a l  
}  
  
#Perform GridSearchCV  
grid_search_rf = GridSearchCV(rf, param_grid_rf, cv = cv, n_jobs = -1, scoring = 'a  
grid_search_rf.fit(X_train, y_train)  
best_rf = grid_search_rf.best_estimator_  
best_rf_params = grid_search_rf.best_params_  
print("RF best parameters: ", best_rf_params)
```

RF best parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 600}

3. Support Vector Machine model

```
In [32]: #Define hyperparameter grids  
param_grid_svm = {  
    'C': [800, 900, 1000, 1100], #Regularization strength; inversely proportional t  
    'gamma': ['scale', 'auto'], #Kernel coefficient for 'rbf', 'poly' and 'sigmoid'  
    'kernel': ['linear', 'poly', 'rbf', 'sigmoid'], #Kernel type to be used  
}  
  
#Perform GridSearchCV  
grid_search_svm = GridSearchCV(svm, param_grid_svm, cv = cv, n_jobs = -1, scoring =  
grid_search_svm.fit(X_train, y_train)  
best_svm = grid_search_svm.best_estimator_  
best_svm_params = grid_search_svm.best_params_  
print("SVM best parameters: ", best_svm_params)
```

SVM best parameters: {'C': 900, 'gamma': 'auto', 'kernel': 'rbf'}

4. K Nearest Neighbours Model

```
In [33]: #Define hyperparameter grids  
param_grid_knn = {  
    'n_neighbors': [1, 2, 3], #Number of default neighbors  
    'weights': ['uniform', 'distance'], #Weight function used in prediction  
    'metric': ['euclidean', 'manhattan', 'minkowski'], #Distance metric to use for  
}
```

```
#Perform GridSearchCV
grid_search_knn = GridSearchCV(knn, param_grid_knn, cv = cv, n_jobs = -1, scoring =
grid_search_knn.fit(X_train, y_train)
best_knn = grid_search_knn.best_estimator_
best_knn_params = grid_search_knn.best_params_
print("KNN best parameters: ", best_knn_params)
```

KNN best parameters: {'metric': 'euclidean', 'n_neighbors': 1, 'weights': 'uniform'}

5. XGBoost Classifier Model

```
In [34]: #Define hyperparameter grids
param_grid_xgb = {
    'n_estimators': [60, 70, 80, 90, 100], # Number of trees
    'learning_rate': [0.01, 0.1, 0.5, 1], # Learning rate
    'max_depth': [7, 8, 9, 10], # Maximum depth of the tree
    'subsample': [0.5, 0.6, 0.7, 0.8], # Subsample ratio of the training instance
    'colsample_bytree': [0.8, 0.9, 1.0], # Subsample ratio of columns when constr
    'reg_alpha': [0, 0.1], # L1 regularization
    'reg_lambda': [1, 0.1] # L2 regularization
}
```

```
#Perform GridSearchCV
grid_search_xgb = GridSearchCV(xgb, param_grid_xgb, cv = cv, n_jobs = -1, scoring =
grid_search_xgb.fit(X_train, y_train)
best_xgb = grid_search_xgb.best_estimator_
best_xgb_params = grid_search_xgb.best_params_
print("XGB best parameters: ", best_xgb_params)
```

XGB best parameters: {'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 8, 'n_estimators': 70, 'reg_alpha': 0, 'reg_lambda': 0.1, 'subsample': 0.7}

Predictions & Evaluations on train data

```
In [35]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
```

1. Logistic Regression model

```
In [36]: #Predict
y_pred_lr_tr = best_lr.predict(X_train)

#Evaluate
print("Logistic Regression Model (Training Data):")
print("Accuracy:", accuracy_score(y_train, y_pred_lr_tr)*100)
print("Confusion Matrix:\n", confusion_matrix(y_train, y_pred_lr_tr))
```

Logistic Regression Model (Training Data):
Accuracy: 72.42990654205607
Confusion Matrix:
[[275 46]
[131 190]]

2. Random Forest Classifier model

```
In [37]: #Predict
y_pred_rf_tr = best_rf.predict(X_train)

#Evaluate
print("Random Forest Classifier Model (Training Data):")
print("Accuracy:", accuracy_score(y_train, y_pred_rf_tr)*100)
print("Confusion Matrix:\n", confusion_matrix(y_train, y_pred_rf_tr))
```

Random Forest Classifier Model (Training Data):
Accuracy: 100.0
Confusion Matrix:
[[321 0]
[0 321]]

3. Support Vector Machine model

```
In [38]: #Predict
y_pred_svm_tr = best_svm.predict(X_train)

#Evaluate
print("SVM Model (Training Data):")
print("Accuracy:", accuracy_score(y_train, y_pred_svm_tr)*100)
print("Confusion Matrix:\n", confusion_matrix(y_train, y_pred_svm_tr))
```

SVM Model (Training Data):
Accuracy: 94.70404984423676
Confusion Matrix:
[[319 2]
[32 289]]

4. K Nearest Neighbours Model

```
In [39]: #Predict
y_pred_knn_tr = best_knn.predict(X_train)

#Evaluate
print("KNN Model (Training Data):")
print("Accuracy:", accuracy_score(y_train, y_pred_knn_tr)*100)
print("Confusion Matrix:\n", confusion_matrix(y_train, y_pred_knn_tr))
```

KNN Model (Training Data):
Accuracy: 100.0
Confusion Matrix:
[[321 0]
[0 321]]

5. XGBoost Classifier Model

```
In [40]: #Predict
y_pred_xgb_tr = best_xgb.predict(X_train)

#Evaluate
print("XGBoost Classifier Model (Training Data):")
print("Accuracy:", accuracy_score(y_train, y_pred_xgb_tr)*100)
print("Confusion Matrix:\n", confusion_matrix(y_train, y_pred_xgb_tr))
```

XGBoost Classifier Model (Training Data):

Accuracy: 100.0

Confusion Matrix:

```
[[321  0]
 [ 0 321]]
```

Predictions & Evaluations on test data

1. Logistic Regression model

```
In [41]: #Predict
y_pred_lr = best_lr.predict(X_test)

#Evaluate
print("Logistic Regression Model (Testing Data):")
print("Accuracy:", accuracy_score(y_test, y_pred_lr)*100)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_lr))
print("Classification Report:", classification_report(y_test, y_pred_lr))
```

Logistic Regression Model (Testing Data):

Accuracy: 62.28070175438597

Confusion Matrix:

```
[[22  7]
 [36 49]]
```

Classification Report:

		precision	recall	f1-score	support
--	--	-----------	--------	----------	---------

0	0.38	0.76	0.51	29
---	------	------	------	----

1	0.88	0.58	0.70	85
---	------	------	------	----

accuracy		0.62	114
----------	--	------	-----

macro avg	0.63	0.67	0.60	114
-----------	------	------	------	-----

weighted avg	0.75	0.62	0.65	114
--------------	------	------	------	-----

2. Random Forest Classifier model

```
In [42]: #Predict
y_pred_rf = best_rf.predict(X_test)

#Evaluate
print("Random Forest Classifier model (Testing Data):")
print("Accuracy:", accuracy_score(y_test, y_pred_rf)*100)
```

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
print("Classification Report:", classification_report(y_test, y_pred_rf))
```

Random Forest Classifier model (Testing Data):

Accuracy: 72.80701754385966

Confusion Matrix:

```
[[17 12]
```

```
[19 66]]
```

Classification Report:

		precision	recall	f1-score	support
	0	0.47	0.59	0.52	29
	1	0.85	0.78	0.81	85
	accuracy			0.73	114
	macro avg	0.66	0.68	0.67	114
	weighted avg	0.75	0.73	0.74	114

3. Support Vector Machine model

```
In [43]: #Predict
y_pred_svm = best_svm.predict(X_test)

#Evaluate
print("Support Vector Machine model (Testing Data):")
print("Accuracy:", accuracy_score(y_test, y_pred_svm)*100)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
print("Classification Report:", classification_report(y_test, y_pred_svm))
```

Support Vector Machine model (Testing Data):

Accuracy: 68.42105263157895

Confusion Matrix:

```
[[17 12]
```

```
[24 61]]
```

Classification Report:

		precision	recall	f1-score	support
	0	0.41	0.59	0.49	29
	1	0.84	0.72	0.77	85
	accuracy			0.68	114
	macro avg	0.63	0.65	0.63	114
	weighted avg	0.73	0.68	0.70	114

4. K Nearest Neighbours Model

```
In [44]: #Predict
y_pred_knn = best_knn.predict(X_test)

#Evaluate
print("K Nearest Neighbours Model (Testing Data):")
print("Accuracy:", accuracy_score(y_test, y_pred_knn)*100)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
print("Classification Report:", classification_report(y_test, y_pred_knn))
```

K Nearest Neighbours Model (Testing Data):

Accuracy: 66.66666666666666

Confusion Matrix:

```
[[15 14]
```

```
[24 61]]
```

Classification Report:

		precision	recall	f1-score	support
	0	0.38	0.52	0.44	29
	1	0.81	0.72	0.76	85
	accuracy		0.67		114
	macro avg	0.60	0.62	0.60	114
	weighted avg	0.70	0.67	0.68	114

5. XGBoost Classifier Model

```
In [45]: #Predict
y_pred_xgb = best_xgb.predict(X_test)

#Evaluate
print("XGBoost Classifier Model (Testing Data):")
print("Accuracy:", accuracy_score(y_test, y_pred_xgb)*100)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_xgb))
print("Classification Report:", classification_report(y_test, y_pred_xgb))
```

XGBoost Classifier Model (Testing Data):

Accuracy: 72.80701754385966

Confusion Matrix:

```
[[16 13]
```

```
[18 67]]
```

Classification Report:

		precision	recall	f1-score	support
	0	0.47	0.55	0.51	29
	1	0.84	0.79	0.81	85
	accuracy		0.73		114
	macro avg	0.65	0.67	0.66	114
	weighted avg	0.74	0.73	0.73	114

Conclusion

```
In [46]: '''
According to the highest accuracy and other evaluation metrics like precision, recall,
F1-score, the most suitable models would be RF and XGBoost. (with around 72% of accuracy)
'''
```

```
Out[46]: '\nAccording to the highest accuracy and other evaluation metrics like precision,
recall and \nF1-score, the most suitable models would be RF and XGBoost. (with around 72% of accuracy)\n'
```

Combine predictions of RF and XGBoost models using an ensemble method - Voting Classifier

```
In [47]: from sklearn.ensemble import VotingClassifier

estimators = [('rf_clf', best_rf), ('xgb', best_xgb)]
voting_clf = VotingClassifier(estimators = estimators, voting = 'soft')
#"Soft" - For voting in averaging probabilities
voting_clf.fit(X_train, y_train)

#Make predictions
y_pred_voting_clf = voting_clf.predict(X_test)

#Evaluate the ensemble model
print("Voting Classifier Model:")
print("Accuracy:", accuracy_score(y_test, y_pred_voting_clf) * 100)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_voting_clf))
print("Classification Report:\n", classification_report(y_test, y_pred_voting_clf))
```

Voting Classifier Model:

Accuracy: 73.68421052631578

Confusion Matrix:

[[16 13]

[17 68]]

Classification Report:

	precision	recall	f1-score	support
0	0.48	0.55	0.52	29
1	0.84	0.80	0.82	85
accuracy			0.74	114
macro avg	0.66	0.68	0.67	114
weighted avg	0.75	0.74	0.74	114

Predict a new observation (whether an individual is a patient or not)

```
In [48]: import numpy as np

feature_names = ['Age', 'Gender', 'Total_Bilirubin', 'Alkaline_Phosphotase',
                 'Alamine_Aminotransferase', 'Aspartate_Aminotransferase', 'Total_P',
                 'Albumin', 'Albumin_and_Globulin_Ratio']
new_obs = np.array([[0.5, 0, 0.3, 0.1, -0.1, 0.2, 0.3, 0.2, 0.1]])

new_obs_data = pd.DataFrame(new_obs, columns = feature_names)

#Predicting using RF and XGBoost models seperately
pred_rf = best_rf.predict(new_obs_data)
print("RF Prediction: ", pred_rf[0])
pred_xgb = best_xgb.predict(new_obs_data)
print("XGB Prediction: ", pred_xgb[0])
```

```
#Predicting using ensemble model  
predict_voting = voting_clf.predict(new_obs_data)  
print("Ensemble Model Prediction ", predict_voting[0])
```

RF Prediction: 1

XGB Prediction: 1

Ensemble Model Prediction 1

The combined model predicts that the individual with given diagnosis is a patient.